



Deep Learning vs Traditional Quant Strategies: A Comparative Study Using CNN-LSTM for Portfolio Optimization

AMOD-5310H: Quantamental Investing and AI

Group 9

Quang Nguyen – ID: 843065

Junyi Chen – ID: 0861615

Yu Zhang – ID: 0862920

Instructor: Prof. Nicholas Faulkner

Table of Contents

Executive Summary	Error! Bookmark not defined.
1. Introduction	5
2. Methodology	6
2.1 Summary of Traditional Strategies (Assignments 1–3).....	6
2.1.1 MPT Portfolio	6
2.1.2 Quantile Regression.....	6
2.1.3 Fama-French Five-Factor Model.....	6
2.1.4 Hybrid Strategy (Signal Processing & Regime Detection).....	6
2.2 LSTM-Based Strategy	8
2.2.1 Motivation for Using LSTM	8
2.2.2 Model Architecture.....	8
2.2.3 Feature Engineering.....	9
2.2.4 Training Process	10
2.2.5 Prediction and Portfolio Construction	11
2.3 Risk Measures Applied	11
3. Data and Implementation.....	13
3.1 Data Collection.....	13
3.2 Preprocessing	13
3.3 Implementation Overview	13
4. Results and Analysis	14
4.1 Return and Risk Comparison	14
4.1.1 Overview of Key Metrics	14
4.1.2 Visual Insights	15
4.1.3 Statistical Significance Testing of Performance Differences	17
4.2 Signal Stability and Turnover.....	18
4.3 Computational Complexity	18
4.4 Trade-off Discussion.....	19
5. Conclusion	20
6. References	22
7. Appendices	22

Table of Figures

Figure 1.	Workflow of the Hybrid Strategy	8
Figure 2.	CNN-LSTM architecture with self-attention for Weekly Return Prediction.	9
Figure 3.	Quadrant-Based Multi-Metric Trade-Off Analysis of Portfolio Strategies	16
Figure 4.	Cumulative Return Comparison between Smart Beta Portfolios	16
Figure 5.	Cumulative Return Comparison between Smart Alpha Portfolios	17
Figure 6.	CNN-LSTM Mean-Variance Portfolio Weights Over Time	23
Figure 7.	CNN-LSTM Hybrid Portfolio Weights Over Time	23
Figure 8.	CNN-LSTM Signal Portfolio Weights Over Time	23
Figure 9.	CNN-LSTM Regime Portfolio Weights Over Time	24
Figure 10.	Heatmap of paired t-test p-values for monthly excess return differences between portfolio strategies.	24
Figure 11.	Significance matrix for paired t-tests on monthly excess returns between all portfolio strategies.	25

List of Tables

Table 1.	Summary of methods used across all assignments	5
Table 2.	Technical indicators	7
Table 3.	Reproducibility Hyperparameter Table	11
Table 4.	Performance Comparison of Portfolio Strategies	15
Table 5.	Average Turnover across Strategies	18
Table 6.	The Use of Generative AI	22

Executive Summary

Introduction:

This project evaluated **traditional quantitative, smart beta, and deep learning–based** portfolio strategies in the Canadian equity market. It builds on methods from Assignments 1–3, extending them with an **LSTM-based predictive model** integrated into portfolio optimization.

Objective:

To compare the performance of diverse portfolio construction methods in terms of **returns, risk measures** (variance, CVaR, Risk Parity Deviation), and **diversification**, and to assess the **trade-offs** between model complexity, interpretability, and performance.

Methodology:

- **Data:** Weekly adjusted prices for 12 TSX-listed large-cap stocks (Aug 2017 – Jun 2025), evaluation from Jan 2020 onward.
- **Strategies Tested:** Modern Portfolio Theory (MPT), Quantile Regression, Fama–French Five-Factor, Hybrid (Regime + Signal), Equal Weight, and four CNN-LSTM variants with self-attention.
- **Evaluation Metrics:** Return (CAGR, Sharpe ratio, total return), risk (variance, CVaR at 95%), and diversification (Risk Parity Deviation)..
- **Statistical Testing:** Paired t-tests on monthly excess returns to assess the significance of performance differences.

Key Findings:

- **Top Risk-Adjusted:** Hybrid (Regime + Signal) - Sharpe **1.46**, CAGR **0.79**, moderate RPD (**1.55**).
- **Top Absolute Return:** MPT Optimized - CAGR **0.85**, but highest volatility and worst tail risk (CVaR **–0.2102**).
- **Best Downside Protection:** Signal Only - CVaR **–0.0895**.
- **Deep Learning Performance:** CNN-LSTM Hybrid Score outperformed other ML variants (Sharpe **1.32**), while CNN-LSTM Mean-Variance underperformed.
- **Most Diversified:** Equal Weight — lowest RPD (**0.60**).
- **Statistical Testing:** Only a small subset of differences, mainly involving Equal Weight versus higher-performing models, were statistically significant at the 5% level.

Recommendations:

- **Retail/Passive Investors:** Equal Weight or Quantile Regression for simplicity and stable diversification.
- **Return-Seeking Investors:** MPT Optimized for higher capital growth, accepting higher risk.
- **Stability-Focused Investors:** Signal Only or Regime Only for adaptive downside protection.
- **Institutional/Advanced Users:** Hybrid (Regime + Signal) or CNN-LSTM Hybrid Score for the best balance of adaptability and performance, provided sufficient computational resources are available.

Conclusion:

Hybrid strategies that **integrate machine learning forecasts with regime and signal overlays** deliver **adaptive, resilient, and well-balanced** portfolios. While deep learning offers performance gains, its benefits are most pronounced when combined with domain-specific signals, and not all improvements are statistically significant. For many investors, simpler approaches may remain more practical given their lower complexity and cost.

1. Introduction

1.1 Project Objectives

This project aims to synthesize traditional quantitative investment strategies with advanced deep learning techniques to construct and evaluate dynamic portfolio allocation models. Specifically, we implement and test a **CNN-LSTM-based prediction strategy** and compare its performance against established methods developed in Assignments 1 through 3. The core objective is to determine whether modern machine learning models can offer measurable improvements in return forecasting and portfolio optimization compared to traditional statistical approaches.

Our comparative analysis focuses on the following dimensions:

- **Return performance**, measured through cumulative and annualized returns;
- **Risk metrics**, including variance, Conditional Value-at-Risk (CVaR), and risk parity;
- **Model complexity and interpretability**, weighing the trade-offs between advanced predictive capabilities and implementation feasibility.

1.2 Motivation & Relevance

In recent years, the integration of **artificial intelligence (AI)** and **deep learning** into financial modeling has gained significant attention. Recurrent neural networks, particularly **Long Short-Term Memory (LSTM)** models, have shown strong potential in capturing complex, nonlinear, and temporal relationships in financial time series. These models are particularly well-suited to address the limitations of traditional approaches, which often rely on static assumptions and linear frameworks.

Evaluating whether LSTM-based strategies can outperform traditional models such as Modern Portfolio Theory or factor-based strategies is essential for informing both academic research and practical asset management. If machine learning models can deliver superior performance or risk control, they may justify the additional computational cost and complexity they entail.

1.3 Overview of Strategies

This project builds upon the foundational work in Assignments 1–3, which explored a range of classical portfolio construction techniques:

Strategy	Description
Modern Portfolio Theory (MPT)	A classical mean-variance optimization framework that seeks to maximize expected return for a given level of risk based on historical covariances.
Quantile Regression	A robust regression method used to forecast conditional quantiles of returns, enabling momentum-based portfolio construction across different deciles.
Fama-French Five-Factor Model	A factor-based approach that explains asset returns through exposures to market, size, value, profitability, and investment style factors.
Hybrid (MACD-HMM Regime-Switching Model)	A hybrid technique combining MACD indicators for trend detection with Hidden Markov Models to probabilistically detect market regimes.
CNN-LSTM-Based Prediction Strategy	A deep learning approach leveraging convolutional layers for local pattern detection and LSTM layers for temporal dependencies to forecast returns.

Table 1. Summary of methods used across all assignments

Table 1 shows the strategies that will be evaluated side-by-side under consistent market conditions and data periods, using common performance and risk metrics to ensure a fair and robust comparison.

2. Methodology

2.1 Summary of Traditional Strategies (Assignments 1–3)

2.1.1 MPT Portfolio

Modern Portfolio Theory (MPT) Portfolio was constructed based on the principle of mean-variance optimization, aiming to maximize portfolio utility by balancing expected return and risk. In our optimization, the objective function was:

$$\max_{\omega} (\mu^T \omega - \lambda \omega^T \Sigma \omega)$$

Where μ is the expected return vector estimated using an exponentially weighted moving average, Σ is the covariance matrix of returns, and λ is the risk aversion coefficient. The portfolio was constrained to fully invest (weights sum to 1), long-only (each weight ≥ 0), and **monthly turnover limit to 30%** relative to the previous period. The optimization process was performed using a 12-month rolling window. Out-of-sample portfolio returns were calculated using next-month realized returns, while in-sample Sharpe ratios were recorded to assess risk-adjusted performance.

2.1.2 Quantile Regression

The Quantile Regression strategy forecasts was implemented to forecast return distributions rather than just the mean. For each stock, a 6-month rolling window regression were fitted for the 90th and 10th percentiles of return against time:

$$\text{Quantile}_q(R_{i,t}) = \alpha_q + \beta_q \cdot t$$

The beta spread ($\beta_{90} - \beta_{10}$) was calculated and then served as a momentum-like signal that reflects the asymmetry of return trends. To compare across stocks, beta spreads were standardized into z-scores. The top 5 ranked stocks were chosen to construct a z-score weighted portfolio, and the remaining stocks were initially weighted at zero. To mitigate transaction cost, we applied a soft turnover constraint of 20%. For example, if the overall monthly change in weights exceeded the 20% threshold, each stock's weight changes will only be limited to a max fraction ($\frac{20\%}{\text{Raw turnover}}$) over the targeted weight change. For comparison, an OLS-based momentum portfolio and an equal-weighted benchmark were also evaluated.

2.1.3 Fama-French Five-Factor Model

In Fama-French Five-Factor model, expected returns were derived from exposures to the five systematic risk factors: market excess return (MKT), size (SMB), value (HML), profitability (RMW), and investment (CMA). For each stock, we estimated factor loadings (betas) using historical excess returns (i.e. return minus risk-free rate):

$$R_{i,t} - R_{f,t} = \alpha_i + \beta_{MKT}MKT_t + \beta_{SMB}SMB_t + \beta_{HML}HML_t + \beta_{RMW}RMW_t + \beta_{CMA}CMA_t + \varepsilon_t$$

These estimated betas were then used to estimate each stock's expected return for the next period, and then the covariance matrix was calculated accordingly. The portfolio was optimized monthly by maximizing the Sharpe ratio under the following constraints:

- Long-only weights
- Full investment ($\sum \text{weights} = 1$)
- Turnover limit less than 20%

Using updated factor loadings and expected returns, the optimizer tilted allocations toward stocks with a strong multi-factor profile.

2.1.4 Hybrid Strategy (Signal Processing & Regime Detection)

The hybrid strategy integrates technical indicators with regime-switching theory to guide portfolio optimization. Each week, a discrete market signal - bullish (+1), bearish (−1), or neutral (0) - is generated from three indicators: MACD histogram smoothing, Relative Strength Index (RSI), and volatility filtering (**Table 2**). A bullish signal occurs when the smoothed MACD histogram is positive, RSI exceeds 50, and portfolio volatility is below the 70th percentile of its historical range. A bearish signal occurs when the

MACD histogram is negative, RSI falls below 50, and volatility remains low. Signals are smoothed over a three-week window to filter short-term noise.

Indicator	Description
MACD histogram	Calculated from the portfolio's average return using the difference between the 12-week and 26-week exponential moving averages, further smoothed by subtracting a 9-week EMA signal line. A positive histogram implies upward momentum, while a negative value suggests weakening trends.
RSI	The average gains to average losses over a 14-week rolling window, transformed into [0,100) range.
Volatility	5-week rolling standard deviation of portfolio returns.

Table 2. *Technical indicators*

A two-state Gaussian Hidden Markov Model (HMM) is trained using standardized MACD and MACD signal line. It classifies each week into **risk-on (regime 1)** or **risk-off (regime 0)**. Regime 1 is associated with sustained upward trends, lower volatility, and positive return expectations, prompting greater risk-taking. Regime 0 corresponds to weakening momentum, higher volatility, and increased drawdown risk, leading investors to reduce exposure.

Expected returns are estimated weekly via regression on market and momentum factors. Portfolio optimization then adapts to the regime: maximizing the Sharpe ratio in risk-on states and minimizing Conditional Value at Risk (CVaR) in risk-off states. Both objectives are subject to four constraints:

1. Long-only weights
2. Full investment (sum of weights equals 1)
3. Turnover below 20%
4. Minimum momentum exposure depending on the signal (Bullish ≥ 0.5 , Neutral ≥ 0.25 , Bearish ≥ 0)

This framework enables the portfolio to shift dynamically, adopting a defensive stance in adverse conditions and a more aggressive allocation when market conditions are favorable. For performance evaluation, the **Hybrid** strategy was implemented alongside the **Signal-only** and **Regime-only** strategies to assess the incremental benefits of combining both approaches.

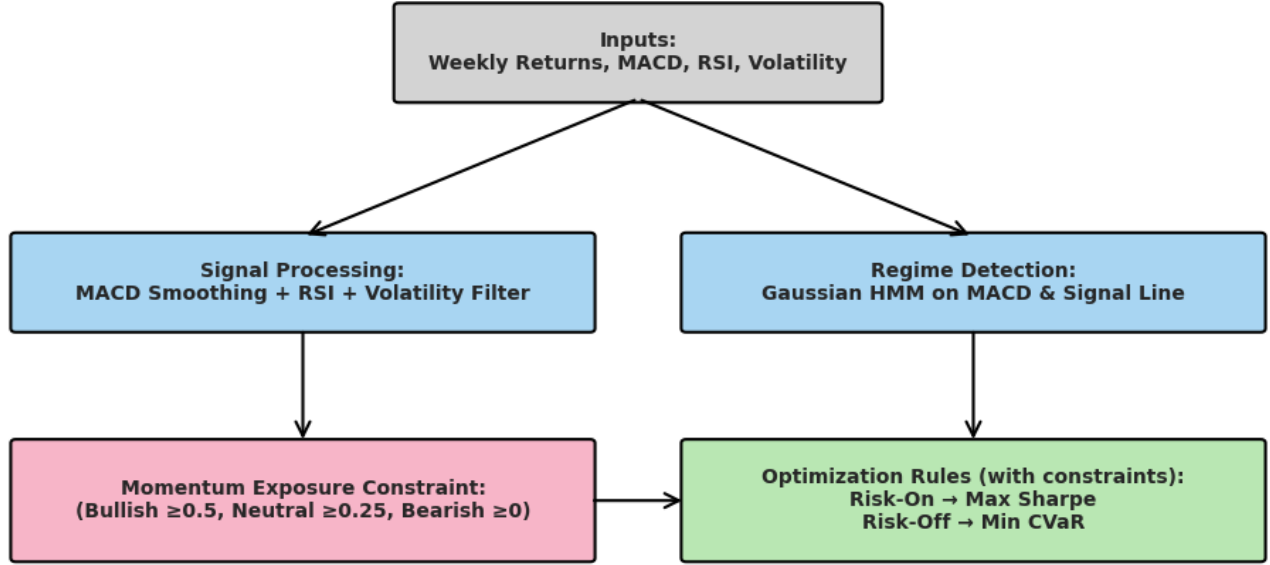


Figure 1. Workflow of the Hybrid Strategy

2.2 LSTM-Based Strategy

2.2.1 Motivation for Using LSTM

Recurrent Neural Networks (RNNs), particularly Long Short-Term Memory (LSTM) models, effectively model financial time series due to their capability to capture temporal dependencies and non-linear patterns (Lai et al., 2018). Traditional financial models often rely on fixed assumptions or static parameters, limiting their adaptability in dynamic market environments. LSTMs, however, are designed to retain relevant past information and selectively forget irrelevant data, enabling robust modeling of noisy, autocorrelated financial returns. Fischer and Krauss (2018) demonstrated that an LSTM-based portfolio achieved significant out-of-sample returns, outperforming traditional forecasting methods, highlighting the LSTM's suitability for predicting stock market behavior.

Convolutional Neural Networks (CNNs) have also shown considerable effectiveness in extracting short-term, local patterns within sequential financial data. Zeng et al. (2023) showed that combining CNN and transformer architectures improved the prediction accuracy by effectively capturing short-term market movements and reducing signal noise. Integrating CNN layers prior to LSTM layers can thus enhance predictive performance by extracting meaningful local features, making it easier for subsequent recurrent layers to model longer-term dependencies (Lai et al., 2018).

Self-attention mechanisms further enhance CNN-LSTM models by dynamically emphasizing critical time steps in input sequences. Vaswani et al. (2017) introduced attention mechanisms that improved the predictive accuracy of sequential models by effectively identifying and weighting important periods within the data. Similarly, Qin et al. (2017) found that incorporating attention mechanisms in recurrent networks significantly enhanced prediction performance on various time series tasks, emphasizing the attention mechanism's role in capturing global temporal dependencies and improving model interpretability.

Thus, integrating CNN, LSTM, and self-attention components creates a powerful framework that effectively combines local pattern extraction, long-term dependency modeling, and dynamic information weighting, leading to improved financial time series forecasting performance.

2.2.2 Model Architecture

To capture the complex temporal patterns in financial returns, we designed a deep sequence learning model that combines convolutional, recurrent, and attention-based components. A one-dimensional

convolutional layer (Conv1D) first extracts short-term patterns and smooths local fluctuations, while batch normalization improves training stability and max pooling reduces dimensionality for improved efficiency.

An LSTM layer then models longer-term dependencies, which is essential in financial markets where trends and cycles can span multiple weeks. A self-attention mechanism follows, allowing the model to focus on the most informative periods and reduce the influence of noise, a common challenge in financial data. The output is processed through two dense layers with ReLU activation and dropout, balancing predictive power with overfitting control.

The model predicts the next week's return using a Huber loss function, which is robust to market outliers, and is optimized with Adam for fast, stable convergence. This architecture not only captures both short- and long-term market dynamics but also improves interpretability through attention weights, helping identify which historical patterns are most influential in the forecast. The complete design is illustrated in **Figure 2**.

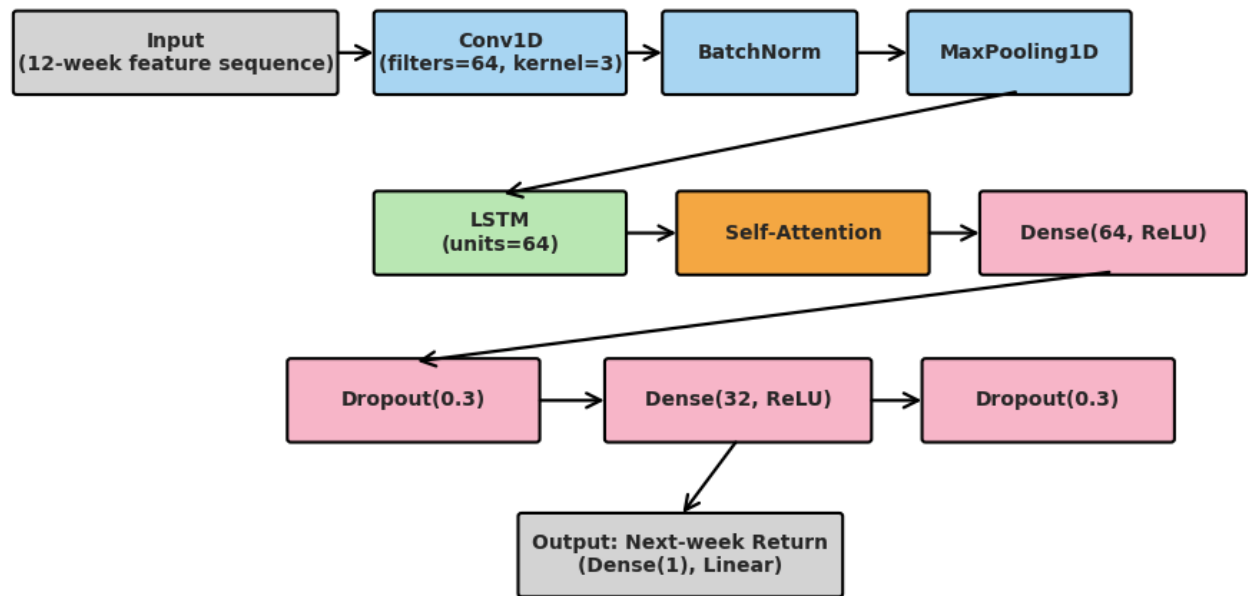


Figure 2. CNN-LSTM architecture with self-attention for Weekly Return Prediction.

2.2.3 Feature Engineering

To enhance predictive performance, we constructed a set of technical indicators derived directly from the weekly return series instead of raw price data. This approach allows the model to learn from more informative features that reflect underlying market behavior, such as trend momentum, volatility, and potential reversals. The first set of indicators includes the Moving Average Convergence Divergence (MACD), its signal line, and the MACD histogram, all computed using exponential moving averages of return series. These indicators are widely used in practice to detect momentum shifts.

We also compute the Relative Strength Index (RSI), which serves as a bounded momentum oscillator capturing overbought or oversold market conditions based on recent return strength. To measure risk, we include rolling volatility, defined as the standard deviation of returns over a 12-week window. Additionally, Bollinger Bands are calculated using a 20-week moving average and standard deviation, providing dynamic upper and lower bounds that signal breakout or mean-reversion opportunities.

Lagged returns from one, two, and three weeks prior are also included as features, allowing the model to directly observe recent return behavior. Finally, we add the lagged weekly return of the S&P/TSX Composite Index (SPTSX) to capture market-wide effects and systematic risk exposure. All features are normalized using MinMaxScaler to ensure that their magnitudes are consistent and suitable for training a neural network. This standardization also aids in convergence and helps prevent dominance of any particular input variable due to scale differences.

2.2.4 Training Process

To evaluate model performance in a way that mimics real-world trading conditions and avoids look-ahead bias, we implemented a **walk-forward training approach**. This method simulates how the model would operate in practice: it is trained on the most recent historical data, used to make predictions for a future period, and then retrained as new data becomes available.

In our setup, the model is retrained every few weeks using a **rolling training window of 104 weeks (approximately 2 years)**. After each retraining, the model forecasts returns for the **next 13 weeks (a quarter)**, which serve as the out-of-sample test period. This process repeats every 13 weeks throughout the dataset, creating a series of rolling forecasts and allowing the model to continuously adapt to changing market conditions.

Each training sample fed into the model consists of a **12-week sequence of historical features**, including technical indicators and lagged returns. This sequence length is designed to give the LSTM and attention layers enough context to learn temporal patterns in the data. After each training iteration, the model produces predictions for the immediate future, and these predicted returns are collected across all stocks for use in the portfolio optimization stage.

To evaluate the quality of the predictions, we compare the model’s output to the actual observed returns using two common performance metrics: the **coefficient of determination (R^2)** and **root mean squared error (RMSE)**. These metrics provide insight into how well the model captures return dynamics and the magnitude of its forecast errors. This process is repeated for each stock individually across the full dataset, and the resulting predicted return matrix is then used in the optimization stage to construct a dynamically rebalanced portfolio strategy. To facilitate reproducibility, the key training and model hyperparameters are summarized in **Table 3**.

Hyperparameter	Value	Description
Sequence length	12 weeks	Number of past weeks of features used as input.
Prediction horizon	13 week	Forecast target is next 13 weeks’ return.
Training window size	104 weeks	Rolling lookback period for walk-forward retraining.
Test window size	13 weeks	Out-of-sample evaluation period after each retraining.
Conv1D filters	64	Number of convolutional filters.
Conv1D kernel size	3	Size of the convolutional kernel (in time steps).
LSTM units	64	Number of units in the LSTM layer.
Dense layers	[64 (ReLU), 32 (ReLU)]	Fully connected layers after attention.
Dropout rate	0.3	The fraction of neurons dropped to reduce overfitting.
Attention mechanism	Scaled dot-product self-attention	Highlights important time steps.
Loss function	Huber loss	Balances MAE and MSE, robust to outliers.
Optimizer	Adam	Adaptive learning rate optimizer.

Batch size	128	Number of samples per gradient update.
Epochs per retraining	50	Maximum number of training passes over data.
Feature scaling	MinMaxScaler (0–1)	Normalizes input features for stability.
Turnover limit	30%	Maximum weekly change in portfolio weights.

Table 3. Reproducibility Hyperparameter Table

2.2.5 Prediction and Portfolio Construction

The predicted weekly stock returns from the CNN-LSTM model were used to construct four different portfolio strategies. Each one combined the model’s predictive power with practical investment constraints and techniques to adapt to different market conditions.

The first strategy was the **CNN-LSTM Mean-Variance Portfolio**. In this approach, the model’s predicted returns were treated as the expected return vector in a classical mean-variance optimization. The covariance matrix was calculated using a 12-week trailing window of historical returns. To ensure the strategy remained practical and avoided excessive trading, a turnover constraint limited the weekly change in portfolio weights to 20%. The portfolio was rebalanced weekly using updated predictions and risk estimates, allowing it to adjust dynamically as market conditions changed.

The second strategy was the **CNN-LSTM Signal-Based Portfolio**. This method added a technical signal component by analyzing the trend in predicted returns over the past four weeks. A linear regression model estimated the slope of the return trend, which was then divided by its volatility to compute a “trend score.” This score indicated whether returns were trending upward or downward, and how strong the trend was. The optimizer used this score to set a target exposure to a macro-style factor (such as sensitivity to a moving average), encouraging the portfolio to tilt toward assets with favorable momentum.

Next, the **CNN-LSTM Regime-Aware Portfolio** used Hidden Markov Models (HMMs) to detect market regimes, such as aggressive (growth) or defensive (risk-off) environments. Each week, the HMM produced probabilities for each regime. If the probability of a defensive regime was higher than 60%, the portfolio focused on minimizing downside risk by optimizing for Conditional Value-at-Risk (CVaR) at the 95% level. If an aggressive regime was more likely, the strategy instead aimed to maximize the Sharpe ratio based on the model’s predictions. This allowed the portfolio to switch between cautious and growth-focused behavior depending on market signals.

Finally, the **CNN-LSTM Hybrid Score Portfolio** combined both regime detection and signal-based exposure control. In defensive regimes, it focused on minimizing CVaR while keeping a neutral stance on the signal. In aggressive regimes, it aimed to maximize the Sharpe ratio, but only if it also aligned with strong positive trend signals. The trend score was used to set a smooth exposure target through a controlled mapping function, helping to balance risk and reward based on both market regime and signal momentum.

Together, these strategies show different ways to turn CNN-LSTM predictions into investment decisions. Each one reflects a different philosophy, whether it’s maximizing returns, managing risk, capturing trends, or adapting to changing markets. By testing and comparing these approaches, we can better understand how deep learning can support smart, data-driven portfolio management.

2.3 Risk Measures Applied

To comprehensively assess the portfolio strategies in this study, we applied three complementary risk measures: **variance**, **Conditional Value-at-Risk (CVaR)**, and **Risk Parity Deviation (RPD)**. Each captures a distinct dimension of portfolio risk, enabling a multidimensional evaluation of performance under varying market conditions.

2.3.1. Variance

Variance is the classical measure of risk in financial theory, quantifying the dispersion of returns around the mean. For a portfolio with the weight vector w and the covariance matrix Σ , variance is computed as:

$$\sigma_p^2 = w^T \Sigma w$$

Higher variance (σ_p^2) indicates greater volatility and uncertainty in portfolio outcomes. In this project, variance served both as a benchmark for return stability and, in the case of mean–variance optimization, as the core objective function. Grounded in Modern Portfolio Theory (MPT), this approach assumes that investors prefer portfolios with lower volatility for a given level of expected return.

2.3.2 Conditional Value-at-Risk (CVaR)

While variance treats upside and downside deviations equally, **CVaR** focuses specifically on extreme downside risk. At the α confidence level (here $\alpha = 95\%$), CVaR measures the expected loss in the worst $1 - \alpha$ fraction of outcomes:

$$CVaR_\alpha = E[R \mid R \leq VaR_\alpha]$$

where VaR_α is the Value-at-Risk at level α , defined as the $(1-\alpha)$ quantile of the return distribution. We estimated CVaR empirically from the weekly return series of each strategy. CVaR is particularly relevant in financial markets, where extreme losses can disproportionately impact long-term performance.

2.3.3 Risk Parity Deviation (RPD)

Portfolio diversification quality was evaluated using **Risk Parity Deviation**, which measures how evenly total portfolio risk is distributed across assets.

First, the **marginal risk contribution** (MRC) of asset i is:

$$MRC_i = \frac{w_i(\Sigma w)_i}{\sigma_p}$$

The **total risk contribution** (TRC) of asset i is:

$$TRC_i = \frac{MRC_i}{\sigma_p}$$

In an ideal risk-parity portfolio, $TRC_i = \frac{1}{N}$ for all i (where N is the number of assets). We define RPD as the root mean square deviation from this ideal:

$$RPD = \sqrt{\frac{1}{N} \sum_i^N (TRC_i - \frac{1}{N})^2}$$

Lower RPD values indicate more balanced risk contributions across assets, while higher values indicate concentration of portfolio risk.

2.3.4 Integrated Risk Framework

Together, these three measures - variance (σ_p^2), $CVaR_{95\%}$, and RPD - provide a robust evaluation framework spanning overall volatility, tail risk, and diversification quality. This multidimensional approach enables more nuanced comparisons between traditional, smart beta, and machine learning–driven strategies.

3. Data and Implementation

3.1 Data Collection

The dataset for this project consists of **weekly adjusted closing prices** of twelve large-cap stocks listed on the **Toronto Stock Exchange (TSX)**. The selected tickers include a diverse mix of sectors such as finance, energy, technology, and healthcare, aiming to provide a balanced representation of the Canadian equity market. The full list of tickers is: ["RY.TO", "MEG.TO", "ATH.TO", "CNQ.TO", "SIA.TO", "MFC.TO", "TECK-B.TO", "BHC.TO", "GIB-A.TO", "CLS.TO", "NPL.TO", "BEP-UN.TO"]. These tickers were chosen based on liquidity, sector representation, and availability of long-term historical data.

To ensure a fair comparison with the portfolio optimization strategies developed in **Assignments 1 through 3**, which were evaluated using data beginning in **January 2020**, we aligned all strategy evaluations to start from that same point. However, because the **CNN-LSTM model** requires substantial training history to make reliable out-of-sample predictions, we collected a total of **eight years of historical data**, from **August 2017 to June 2025**. This extended dataset ensures that the model has sufficient lookback periods for training and validation while preserving the comparability of out-of-sample performance across strategies.

In addition to individual stock data, we retrieved the **weekly adjusted closing prices of the S&P/TSX Composite Index** to serve as both a **market benchmark** and an input feature for modeling. All price data was collected using the **Yahoo Finance API** via the `yfinance` library. Where necessary, we supplemented this with **Bloomberg Terminal** queries using **BQL** for cross-validation or enhanced data granularity.

3.2 Preprocessing

Before modeling, several preprocessing steps were performed to prepare the data for analysis and model input. First, the raw daily prices were resampled into **weekly closing prices**, aligned to **Fridays**, which is a common approach to reduce noise and reflect realistic portfolio rebalancing frequency. From these prices, weekly returns were calculated and used as the primary input for traditional models and for computing technical indicators. After the return calculation, any residual NaNs were dropped to ensure consistency across time steps and assets. The result was a clean, aligned return matrix across all 12 stocks and the TSX index.

To enrich the input features for the deep learning models, we computed a variety of return-based technical indicators, such as MACD, RSI, rolling volatility, and Bollinger Bands, using the `ta` Python package. These indicators help the model detect patterns in momentum, trend strength, and risk, which are not always evident in raw returns. All features were scaled using `MinMaxScaler` before being passed to the LSTM network to ensure numerical stability and efficient convergence during training.

3.3 Implementation Overview

The project was implemented entirely in **Python**, leveraging a range of popular data science and machine learning libraries. For data manipulation and preprocessing, we used **Pandas** and **NumPy**. Technical indicator computation was done using the **ta package**, and in earlier stages, we experimented with `TA-Lib` as an alternative. Deep learning models were built using **TensorFlow** and Keras, while `scikit-learn` was used for scaling, regression metrics, and general utilities.

Each strategy was implemented in modular scripts to support large-scale simulation and retraining. The mean-variance optimization was solved using `scipy.optimize.minimize`, and CVaR was estimated from historical and simulated return distributions. For evaluating risk parity, we computed marginal contributions to portfolio volatility using matrix calculus applied to the covariance matrix.

Walk-forward training and prediction for the **CNN-LSTM model** were executed on a standard CPU environment with GPU T4x2 acceleration for faster training. Due to the manageable size of the dataset and the weekly frequency of training, GPU usage was not mandatory, but it helped reduce model fitting time.

Plots and diagnostic metrics such as R^2 and RMSE were visualized using Matplotlib to evaluate the prediction quality across time.

In summary, the implementation combined traditional financial modeling techniques with modern deep learning approaches, resulting in a robust framework for comparing multiple portfolio strategies under consistent data and execution settings.

4. Results and Analysis

4.1 Return and Risk Comparison

4.1.1 Overview of Key Metrics

To evaluate the effectiveness of the portfolio strategies, we compare both traditional and CNN-LSTM-based models using a consistent set of risk-return metrics. **Table 4** summarizes key indicators, including Compound Annual Growth Rate (CAGR), volatility, Conditional Value-at-Risk (CVaR 95%), Sharpe ratio, maximum drawdown, total return, tracking error, and Risk Parity Deviation (RPD). Traditional strategies like **MPT Optimized** and **Quantile Regression Score** achieve strong returns, with MPT showing the highest **CAGR (0.8522)** and **total return (23.15)**, but also the highest **volatility** and **CVaR**. Hybrid strategies, such as **Hybrid (Regime + Signal)** and **Regime Only**, deliver favorable **Sharpe ratios (~1.43–1.46)** with more balanced risk profiles. Notably, **CNN-LSTM-based portfolios** show competitive performance in terms of return and risk-adjusted measures, especially the **CNN-LSTM Hybrid Portfolio**, which yields a high **Sharpe ratio (1.3241)** and solid **CAGR (0.6896)**. However, most deep learning models exhibit slightly higher **RPD values**, indicating less effective risk diversification. Overall, hybrid and regime-aware strategies outperform equal weighting, while the integration of machine learning adds modest gains with increased complexity.

Strategy	CAGR	Volatility	CVaR 95%	Sharpe Ratio	Max Drawdown	Total Return	Tracking Error	RPD (↓ Better)
Equal Weight	0.2234	0.2527	-0.1409	0.8227	-0.3031	1.83	0.1476	0.60
MPT Optimized	0.8522	0.4989	-0.2102	1.4243	-0.3673	23.15	0.5387	1.73
Quantile Regression Score	0.2995	0.3051	-0.1524	0.9310	-0.3649	2.87	0.2160	1.08
Hybrid (Regime + Signal)	0.7949	0.4494	-0.1211	1.4556	-0.3770	17.03	0.3968	1.55
Signal Only	0.4020	0.2918	-0.0895	1.2025	-0.3331	4.32	0.2305	1.35
Regime Only	0.7568	0.4411	-0.1191	1.4259	-0.3728	15.22	0.3830	1.56
CNN-LSTM Regime Portfolio	0.6556	0.4389	-0.1200	1.2995	-0.3643	10.9800	0.3950	1.80
CNN-LSTM Mean-Var Portfolio	0.4216	0.4611	-0.1255	0.9261	-0.4979	4.6942	0.4046	1.80

CNN-LSTM Signal Portfolio	0.4584	0.4679	-0.1335	0.9951	-0.4793	5.4616	0.4496	1.73
CNN-LSTM Hybrid Portfolio	0.6896	0.4491	-0.1194	1.3241	-0.3657	12.2415	0.4044	1.79

Table 4. Performance Comparison of Portfolio Strategies

4.1.2 Visual Insights


Figure 3 assesses portfolio strategies by comparing **Compound Annual Growth Rate (CAGR)** against three key risk metrics: **Volatility**, **CVaR 95%**, and **Risk Parity Deviation (RPD)**.

In the **CAGR vs Volatility** panel (top-left), the ideal strategies appear in the upper-left quadrant. **MPT Optimized** and **Hybrid (Regime + Signal)** dominate with the highest returns, while maintaining volatility at acceptable levels. **CNN-LSTM Regime** and **CNN-LSTM Hybrid** also perform well, though with slightly more volatility. Simpler models like **Equal Weight** and **Quantile** offer lower returns with lower risk.

In the **CAGR vs CVaR 95%** panel (top-right), **MPT Optimized** achieves the highest CAGR but suffers from the worst tail risk (CVaR = -0.2102). **Hybrid** stands out with strong return and more moderate downside risk. **Signal Only** is the most conservative in terms of CVaR, but yields modest growth.

In the **CAGR vs RPD** panel (bottom), **Equal Weight** exhibits the best diversification (lowest RPD) but the lowest return. **Hybrid** and **Regime-only** achieve strong CAGR with moderately high RPD, suggesting a more concentrated risk. CNN-LSTM models show moderate returns but relatively high RPD, indicating weaker risk parity. This result is confirmed by looking at **Figures 6 to 10** in the Appendix Section.

Overall, the **Hybrid** strategy offers the best return-risk trade-off, the highest Sharpe ratio (1.46), strong CAGR (0.8), controlled CVaR, and lower RPD than all other strategies with CAGR > 0.5. While **MPT Optimized** delivers the highest returns, it consistently incurs the most risk. **CNN-LSTM Regime** and **Hybrid** approaches show promising results, balancing machine learning with regime awareness for improved portfolio performance. However, this comes with elevated **Risk Parity Deviation**, indicating less effective diversification.






Figure 3. Quadrant-Based Multi-Metric Trade-Off Analysis of Portfolio Strategies

Figure 4 presents the cumulative returns of traditional smart beta portfolios and the LSTM-enhanced mean-variance portfolio. The **MPT Optimized** strategy achieved the highest cumulative return over the sample period, particularly after 2023, though it exhibited considerable volatility and drawdowns. In contrast, the **Equal-Weight** and **Quantile Regression Score** strategies delivered modest and steady growth, serving as benchmarks for passive allocation and alternative scoring methods.

The **CNN-LSTM Mean-Variance Portfolio**, which integrates deep learning predictions into a classic optimization framework, consistently outperformed the Equal-Weight and QR Score approaches. Although it did not surpass the MPT Optimized strategy in absolute return, it demonstrated enhanced risk-adjusted performance, validating the potential of deep learning in guiding portfolio construction.




Figure 4. Cumulative Return Comparison between Smart Beta Portfolios

Figure 5 showcases the performance of smart alpha portfolios incorporating regime identification, signal processing, and deep learning. The **Hybrid (Regime + Signal)** strategy outperformed all others, ending the period with the highest wealth index. This confirms the efficacy of combining macroeconomic regimes and momentum-based signals for adaptive allocation.

The **CNN-LSTM Hybrid Score Portfolio** closely tracked the Hybrid strategy's trajectory, suggesting that machine learning adds incremental value when overlaid on a strong structural foundation. Similarly, both the **CNN-LSTM + Regime Portfolio** and **CNN-LSTM Signal Portfolio** delivered robust performance, validating the standalone utility of each approach.

Meanwhile, the **Signal Only** and **Regime Only** portfolios, though simpler in construction, also achieved competitive results, outperforming naive strategies and demonstrating effective downside protection and responsiveness to market shifts.




Figure 5. Cumulative Return Comparison between Smart Alpha Portfolios

The comparative analysis of portfolio strategies reveals that the **Hybrid (Regime + Signal)** portfolio outperformed across most metrics, achieving the highest Sharpe ratio (1.46) and strong CAGR (0.79), highlighting its superior balance of return and risk control. While the **MPT Optimized** strategy delivered the highest total return (23.15), it incurred higher volatility and the worst downside risk (CVaR: -0.21), underscoring the trade-off between return and stability. Among smart beta strategies, the **Signal Only** portfolio exhibited the best downside protection (CVaR: -0.0895), and **Regime Only** achieved strong risk-adjusted performance. Deep learning models further enhanced outcomes when combined with structural signals: the **CNN-LSTM Hybrid Score Portfolio** delivered a Sharpe ratio of 1.32 and robust cumulative returns, validating the effectiveness of hybrid learning-finance integration. In contrast, the standalone **CNN-LSTM Mean-Variance Portfolio** lagged due to higher drawdowns and limited diversification. Overall, the results emphasize that integrating domain-informed signals and regimes into deep learning models yields more adaptive, resilient, and high-performing portfolios.

4.1.3 Statistical Significance Testing of Performance Differences

To complement the descriptive performance analysis, we conducted paired t-tests on monthly excess returns for all strategy pairs to evaluate whether observed differences were statistically significant. This approach compares mean excess returns over overlapping months, controlling for shared market conditions. The resulting significance matrix (Appendix B, Figures 10 and 11) flags pairs with $p < 0.05$ as statistically different at the 95% confidence level.

Findings show that only a **limited number** of comparisons, primarily involving the **Equal Weight portfolio** against **higher-performing models** such as MPT Optimized, Hybrid (Regime + Signal), and certain deep learning hybrids, met the significance threshold. In contrast, **most differences** between non-baseline strategies, including those between MPT, Hybrid, and deep learning-based portfolios, were **not statistically significant**. This suggests that while point estimates of Sharpe ratios and CAGR vary, much

of the observed variation may be attributable to sampling variability over the 65-month evaluation period rather than persistent structural advantages.

4.2 Signal Stability and Turnover

Weight Turnover Comparison

To evaluate how frequently strategies adjust their holdings, we computed the average weekly turnover for each model. As shown in **Table 5**, Traditional models like MPT, Quantile Regression, and Equal Weight exhibited more stable allocations due to longer rebalancing periods or static weights. In contrast, Hybrid (QR and Signal), Signal-only, Regime-only, and CNN-LSTM Signal Portfolio strategies exhibited higher turnover, reflecting their responsiveness to predicted returns.

Strategy	Rebalancing Frequency	Avg Turnover
Equal Weight	Static	0.00
MPT Optimized	Monthly	~0.09 (Monthly)
Quantile Regression Score	Monthly	~0.10 (Monthly)
Hybrid (QR + Signal)	Weekly	~0.13
Signal Only	Weekly	~0.12
Regime Only	Weekly	~0.13
CNN-LSTM Mean-Var Portfolio	Weekly	~0.07
CNN-LSTM Regime Portfolio	Weekly	~0.05
CNN-LSTM Signal Portfolio	Weekly	~0.15
CNN-LSTM Hybrid Score Portfolio	Weekly	~0.04

Table 5. Average Turnover across Strategies

Transaction Cost Implications:

Although dynamic strategies tend to respond faster to market signals, they incur **higher transaction costs** due to frequent weight adjustments. This trade-off must be considered when selecting a portfolio strategy in real-world applications. For instance, **CNN-LSTM Signal Portfolio**, while responsive, showed elevated turnover and higher drawdowns.

4.3 Computational Complexity

4.3.1 Training Times

Traditional portfolio strategies such as Modern Portfolio Theory (MPT), Quantile Regression, and the Fama-French model are highly efficient in terms of computation. These models typically complete execution within seconds to a few minutes on standard hardware, which makes them practical for rapid backtesting and deployment.

In contrast, CNN-LSTM models require significantly more computational resources. This is especially evident under the walk-forward retraining framework used in this study. Each retraining iteration, covering approximately 104 weeks of data, required between 60 and 120 seconds per stock when run on an NVIDIA T4 GPU. When aggregated across all stocks in the portfolio, the total training time reached approximately 10 hours. This highlights the need for access to high-performance computing infrastructure when working with deep learning models in finance.

4.3.2 Prediction Latency

Once trained, the CNN-LSTM model's prediction latency is negligible, less than 1 second per stock, making it suitable for weekly rebalancing cycles.

4.3.3 Model Explainability

Interpretability is a major differentiating factor between traditional and deep learning models.

- MPT uses explicit inputs such as expected returns and covariance matrices, which directly influence portfolio weights.
- The Fama-French model provides clear explanations for asset performance using predefined economic factors such as size and value.

On the other hand, CNN-LSTM models function as complex, nonlinear systems that are more difficult to interpret. Although attention mechanisms can highlight which time periods the model considers important, they do not fully explain the influence of each input feature on the final prediction. This limited transparency introduces a trade-off. While deep learning models may offer improved predictive accuracy, they often lack the interpretability that traditional financial models provide.

4.4 Trade-off Discussion

4.4.1 Complexity vs. Performance

While deep learning models such as CNN-LSTM introduce a high level of computational and implementation complexity, the performance gains they deliver can justify the added cost, particularly when combined with structural signals. For instance, the **CNN-LSTM Hybrid Score Portfolio** outperformed many traditional strategies in terms of the Sharpe ratio and consistency of returns. However, the **CNN-LSTM Mean-Variance Portfolio**, which lacked structural guidance, underperformed, suggesting that deep learning models are not universally superior on their own. The marginal benefit of using LSTM over simpler models depends heavily on how well it is integrated with domain knowledge, such as regime classification or momentum signals. For investors with limited resources or shorter investment horizons, the additional complexity of deep learning may not be worthwhile unless the models are properly augmented and explainable.

4.4.2 Use Case Suitability

The choice of portfolio strategy should be guided by the investor's specific objectives and constraints. For example, **Equal-Weight** and **Quantile Regression Score** portfolios are well-suited for passive or rule-based investors who prioritize simplicity and broad diversification. These strategies require minimal intervention and offer stable, easily replicable allocation methods.

In contrast, **MPT Optimized** portfolios tend to appeal to return-seeking investors who are comfortable taking on higher volatility and tail risk, particularly in bullish market environments. By optimizing for expected return relative to risk, these portfolios can achieve strong capital growth but may expose investors to larger drawdowns during market downturns.

For those seeking a balanced approach, **Signal Only** and **Regime Only** smart beta strategies offer an attractive middle ground. These methods deliver improved risk-adjusted performance while maintaining interpretable logic and a relatively low implementation burden. They adapt allocations based on market signals or regime conditions without requiring complex infrastructure.

Finally, **Hybrid (Regime + Signal)** and **CNN-LSTM Hybrid Score** portfolios are designed for advanced users aiming for robust, adaptive strategies that can navigate varying market regimes. These approaches are particularly well-suited for institutional or data-driven investors who have the resources to maintain and manage the necessary model infrastructure.

Ultimately, there is no one-size-fits-all solution; each strategy serves a different purpose depending on risk tolerance, data availability, and implementation capacity.

5. Conclusion

5.1 Summary of Key Findings:

This project conducted a comprehensive evaluation of portfolio construction methods, comparing traditional, smart beta, and deep learning-enhanced strategies. Among all approaches, the **Hybrid (Regime + Signal)** portfolio achieved the highest Sharpe ratio of 1.46 and a strong Compound Annual Growth Rate (CAGR) of 0.7949. It consistently performed well across all risk–return metrics, including a relatively moderate Risk Parity Deviation (RPD) of 1.55, outperforming both classical optimization and machine learning-only approaches.

The **MPT Optimized** portfolio achieved the highest total return (23.15) and CAGR (0.85) but suffered from elevated volatility, the most negative CVaR (−0.2102), and a high RPD (1.73), indicating concentrated and tail-risk exposure. All **CNN-LSTM-based** strategies (including CNN-LSTM Hybrid, Regime, Signal, and Mean-Variance) exhibited high RPD values (≥ 1.73), signaling poor risk diversification. Nonetheless, the **CNN-LSTM Hybrid** strategy delivered a Sharpe ratio of 1.32, demonstrating that blending deep learning with structural elements such as trend signals and regime filters improves performance compared to using predictions alone.

In contrast, the **CNN-LSTM Mean-Variance** portfolio, which relied solely on forecasted returns, underperformed across all key metrics. **Statistical significance testing** using **paired t-tests** indicated that only a limited number of differences, primarily involving the Equal Weight portfolio versus higher-performing models, were significant at the 5% level. Most other differences were not statistically significant, suggesting that part of the observed variation in Sharpe ratios and returns may be attributable to sampling variability rather than persistent structural advantages. These findings reinforce that integrating financial context into data-driven models yields more effective and balanced portfolio strategies.

Complexity and Justification of LSTM:

The CNN-LSTM models introduced substantial computational complexity and training overhead. However, their performance, particularly in hybrid configurations, supports their use for investors seeking adaptive and data-driven allocation strategies. When paired with structured alpha signals, LSTM-based models can outperform simpler models on a risk-adjusted basis. For investors or institutions with access to computational resources, the added complexity may be justifiable. For others, traditional approaches like MPT or Quantile Regression may remain preferable due to their transparency, ease of implementation, and lower operational cost.

5.2 Computational Constraints:

The CNN-LSTM models introduced substantial computational complexity and training overhead. However, their performance, particularly in hybrid configurations, supports their use for investors seeking adaptive and data-driven allocation strategies. For example, the CNN-LSTM Hybrid Score achieved a Sharpe ratio of 1.32, compared to 1.20 for the Signal Only strategy, an improvement of 0.12. Given the total training time of approximately 10 hours, this translates to roughly 83 minutes of additional compute time per 0.01 Sharpe improvement.

With increased computational capacity, the model could be trained for more epochs or retrained more frequently. This improvement could lead to higher predictive accuracy and better portfolio responsiveness. The CNN-LSTM Mean-Variance strategy, which lacked auxiliary structure, may have been particularly affected by this limitation, as it relied entirely on predictive strength.

5.3 Strategic Recommendations

For retail or passive investors, simple strategies such as Equal-Weight or Quantile Regression are easy to implement and offer modest diversification benefits.

For investors willing to accept higher volatility in pursuit of higher returns, the MPT Optimized portfolio provides strong capital appreciation potential.

For those prioritizing stability and responsiveness, strategies such as Regime Only and Signal Only offer adaptive risk control without the need for deep learning.

For advanced practitioners or institutional investors, hybrid strategies that integrate LSTM-based forecasts with domain-specific signals and regime filters appear to deliver the most favorable balance of return and risk.

5.4 Future Research Directions

While the findings are promising, there remain opportunities for further exploration. Future research could include expanding the asset universe beyond twelve stocks, testing alternative model architectures such as transformers, and evaluating different forecasting horizons. Improvements in explainability techniques and integration of macroeconomic indicators could further enhance model utility. Lastly, scaling the infrastructure to support faster retraining and deeper model tuning could improve out-of-sample performance and portfolio adaptability.

6. References

- Fischer, T., & Krauss, C. (2018). Deep learning with long short-term memory networks for financial market predictions. *European Journal of Operational Research*, 270(2), 654–669. <https://doi.org/10.1016/j.ejor.2017.11.054>
- Lai, G., Chang, W.-C., Yang, Y., & Liu, H. (2018). Modeling long- and short-term temporal patterns with deep neural networks. *Proceedings of the 41st International ACM SIGIR Conference on Research & Development in Information Retrieval (SIGIR '18)*, 95–104. <https://doi.org/10.1145/3209978.3210006>
- Qin, Y., Song, D., Chen, H., Cheng, W., Jiang, G., & Cottrell, G. W. (2017). A dual-stage attention-based recurrent neural network for time series prediction. *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence (IJCAI-17)*, 2627–2633. <https://doi.org/10.24963/ijcai.2017/366>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30, 5998–6008. <https://doi.org/10.48550/arXiv.1706.03762>
- Zeng, Z., Kaur, R., Siddagangappa, S., Rahimi, S., Balch, T., & Veloso, M. (2023). Financial time series forecasting using CNN and Transformer. *arXiv preprint arXiv:2304.04912*. <https://doi.org/10.48550/arXiv.2304.04912>

7. Appendices

A. Use of Generative AI

Category	Details
Tool Used	ChatGPT (OpenAI)
Prompts Given	- Requests for help debugging Python code - Clarifications on portfolio optimization methodologies
Actions Taken	- Generated Python code snippets for CNN-LSTM implementation - Provided guidance on debugging and optimizing financial models
Model Focus	CNN-LSTM for financial time series forecasting
Technical Areas Covered	- Time series feature engineering (MACD, RSI, MA, etc.) - CNN-LSTM architecture - Attention mechanism integration - Rolling training
Optimization Topics	- Mean-variance optimized objective function: Sharpe ratio - Risk parity – CVaR minimization,
Outcome	User received end-to-end support in building, debugging, and evaluating a hybrid deep learning-based portfolio strategy

Table 6. The Use of Generative AI

B. Supplementary Figures or Tables




Figure 6. CNN-LSTM Mean-Variance Portfolio Weights Over Time




Figure 7. CNN-LSTM Hybrid Portfolio Weights Over Time




Figure 8. CNN-LSTM Signal Portfolio Weights Over Time




Figure 9. CNN-LSTM Regime Portfolio Weights Over Time




Figure 10. Heatmap of paired t-test p-values for monthly excess return differences between portfolio strategies.

(Darker colors correspond to smaller p-values. The 0.05 significance threshold is highlighted in the accompanying binary significance matrix.)

Significance Flags (1 = $p < 0.05$)									
EW									
MPT	1								
QR	0	0							
HYB	1	0	1						
SIG	0	0	0	1					
REG	1	0	1	0	0				
LSTM-REG	1	0	0	0	0	0			
LSTM-MV	0	0	0	0	0	0	0		
LSTM-SIG	0	0	0	0	0	0	0	0	
LSTM-HYB	1	0	0	0	0	0	0	0	0
	EW	MPT	QR	HYB	SIG	REG	LSTM-REG	LSTM-MV	LSTM-SIG

Figure 11. Significance matrix for paired *t*-tests on monthly excess returns between all portfolio strategies.

(A value of 1 indicates statistical significance at the 5% level, while 0 indicates no significant difference. The matrix is lower-triangular to display unique pairwise comparisons.)

Abbreviations:

- **EW** – Equal Weight
- **MPT** – MPT Optimized
- **QR** – Quantile Regression Score
- **HYB** – Hybrid (Regime + Signal)
- **SIG** – Signal Only
- **REG** – Regime Only
- **LSTM-REG** – CNN-LSTM Regime Portfolio
- **LSTM-MV** – CNN-LSTM Mean-Variance Portfolio
- **LSTM-SIG** – CNN-LSTM Signal Portfolio
- **LSTM-HYB** – CNN-LSTM Hybrid Score Portfolio

C. Full Code Repository

- GitHub code link: <https://github.com/wangwang2111/canada-quantamental-investing-strategies.git>